

# Artificial Intelligence

## 4 Beyond Classical Search

Russell & Norvig, AI: A Modern Approach, 3rd Ed

### (Local Search)

Lec Zinia Sultana  
CSE Dept, MIST


# Search Heuristics

## Local search techniques and optimization

- ✓ Hill-climbing
- ✓ Simulated annealing
- ✓ Local Beam Search
- ✓ Genetic algorithms
- ✓ Issues with local search

# Local search and optimization


- Previously: systematic exploration of search space.
  - Path to goal is solution to problem
- YET, for some problems path is irrelevant.
  - E.g 8-queens
- Different algorithms can be used
  - Local search



# Local search and optimization

## Applications:

- Integrated-circuit design
- Factory-floor layout
- Job-shop scheduling
- Automatic programming
- Telecommunications
- Network Optimization
- Vehicle Routing




# Local search and optimization

- Local search
  - Keep track of single current state
  - Move only to neighboring states
  - Ignore paths
- Advantages:
  - Use very little memory (**Usually Constant Amount**)
  - Can often find reasonable solutions in large or infinite (continuous) state spaces.

# Local search and optimization

- Local Search algorithms are very useful for solving “Pure Optimization” problems.
- “Pure optimization” problems
  - All states have an objective function (OF)
  - Goal is to find state with max (or min) objective value
  - Does not quite fit into path-cost/goal-state formulation
  - Local search can do quite well on these problems.

# “Landscape” of search



A complete local search algorithm always finds a goal if one exists!!

# Hill-climbing search

**function** HILL-CLIMBING(*problem*) **return** a state that is a local maximum

**input:** *problem*, a problem

**local variables:** *current*, a node.

*neighbor*, a node.

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**loop do**

*neighbor*  $\leftarrow$  a highest valued successor of *current*

**if** VALUE [*neighbor*]  $\leq$  VALUE[*current*] **then return** STATE[*current*]

*current*  $\leftarrow$  *neighbor*




# Hill-climbing (Steepest-Ascent Version) search

- “A loop that continuously moves in the direction of increasing value” – to Uphill
  - terminates when a peak is reached
  - Aka greedy local search
- Value can be either
  - Objective function value
  - Heuristic function value
- Hill climbing does not look ahead of the immediate neighbors of the current state.
- Can randomly choose among the set of best successors, if multiple have the best value
- Characterized as “trying to find the top of Mount Everest in a thick fog while suffering from amnesia”

# Hill climbing and local maxima

- When local maxima exist, hill climbing is suboptimal
- Simple (often effective) solution
  - Multiple random restarts











# Hill-climbing example

- 8-queens problem, complete-state formulation
  - All 8 queens on the board in some configuration
- **Successor function:**
  - move a single queen to another square in the same column.
- Example of a heuristic function  $h(n)$ :
  - the number of pairs of queens that are attacking each other (directly or indirectly)
  - (so we want to minimize this)


# Hill-climbing example

Current state:  $h=17$

Shown is the  $h$ -value for each possible successor in each column



18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14		13	16	13	16
	14	17	15		14	16	16
17		16	18	15		15	
18	14		15	15	14		16
14	14	13	17	12	14	12	18

# A local minimum for 8-queens



A local minimum in the 8-queens state space ( $h=1$ )

# Other drawbacks



- Ridge = sequence of local maxima difficult for greedy algorithms to navigate
- Plateau = an area of the state space where the evaluation function is flat.

# Performance of hill-climbing on 8-queens

- Randomly generated 8-queens starting states...
- 14% the time it solves the problem
- 86% of the time it get stuck at a local minimum
- However...
  - Takes only 4 steps on average when it succeeds
  - And 3 on average when it gets stuck
  - (for a state space with  $\sim 17$  million states)

## Possible solution...sideways moves

- If no downhill (uphill) moves, allow sideways moves in hope that algorithm can escape
  - Need to place a limit on the possible number of sideways moves to avoid infinite loops
- For 8-queens
  - Now allow sideways moves with a limit of 100
  - Raises percentage of problem instances solved from 14 to 94%
- However....
  - 21 steps for every successful solution
  - 64 for each failure



# Hill-climbing variations

- Stochastic hill-climbing
  - Random selection among the uphill moves.
  - The selection probability can vary with the steepness of the uphill move.
  - Usually converges more slowly than steepest-ascent but in some state space, it finds better solutions
- First-choice hill-climbing
  - stochastic hill climbing by generating successors randomly until a better one is found
  - Useful when there are a very large number of successors
- Random-restart hill-climbing
  - Tries to avoid getting stuck in local maxima.

# Hill-climbing with random restarts

- Different variations
  - For each restart: run until termination v. run for a fixed time
  - Run a fixed number of restarts or run indefinitely
- Analysis
  - Say each search has probability  $p$  of success
    - E.g., for 8-queens,  $p = 0.14$  with no sideways moves
  - Expected number of restarts?
  - Expected number of steps taken?

# Expected number of restarts

- Probability of Success =  $p$
- Number of restarts =  $1 / p$
- This means 1 successful iteration after  $(1/p - 1)$  failed iterations
- Let avg. number of steps in a failure iteration =  $f$   
and avg. number of steps in a successful iteration =  $s$
- Therefore, expected number of steps in random-restart hill climbing =  $1 * s + (1/p - 1) f$

So for 8-queens,  $p = 14\%$ ,  $s = 4$ ,  $f = 3$ ,

$$\text{Expected no of moves} = 1 * 4 + (1/0.14 - 1) * 3 = 22$$

With sideways moves,  $p = 94\%$ ,  $s = 21$ ,  $f = 64$

$$\text{Expected no of moves} = 1 * 21 + (1/0.94 - 1) * 64 = 25$$

# Local beam search

- Keep track of  $k$  states instead of one
  - Initially:  $k$  randomly selected states
  - Next: determine all successors of  $k$  states
  - If any of successors is goal  $\rightarrow$  finished
  - Else select  $k$  best from successors and repeat.
- Major difference with random-restart search
  - Information is shared among  $k$  search threads.
- Can suffer from lack of diversity.
  - Stochastic beam search
    - choose  $k$  successors proportional to state quality (randomly).

# Search using Simulated Annealing

- Simulated Annealing = hill-climbing with non-deterministic search (random walk)
- Basic ideas:
  - like hill-climbing identify the quality of the local improvements
  - instead of picking the best move, pick one randomly
  - say the change in objective function is  $d$
  - if  $d$  is positive, then move to that state
  - otherwise:
    - move to this state with probability proportional to  $d$
    - thus: worse moves (very large negative  $d$ ) are executed less often
  - however, there is always a chance of escaping from local maxima

# Physical Interpretation of Simulated Annealing

- A Physical Analogy:
  - Imagine letting a ball roll downhill on the function surface
    - this is like hill-climbing (for minimization)
  - Now imagine shaking the surface, while the ball rolls, gradually reducing the amount of shaking
    - this is like simulated annealing
    - If we shake the surface, we can bounce the ball out of local minimum but hard enough to dislodge from global minimum
- Annealing = physical process used to **temper or harden metals** glass by **heating them to a high temperature** and then **gradually cooling them**, thus allowing the material to reach a **low-energy crystalline state**.

# Simulated annealing

**function** SIMULATED-ANNEALING( *problem*, *schedule*) **return** a solution state

**input:** *problem*, a problem

*schedule*, a mapping from time to temperature

**local variables:** *current*, a node.

*next*, a node.

*T*, a “temperature” controlling the probability of downward steps

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**for** *t*  $\leftarrow$  1 **to**  $\infty$  **do**

*T*  $\leftarrow$  *schedule*[*t*]

**if** *T* = 0 **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  VALUE[*next*] - VALUE[*current*]

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E / T}$

# More Details on Simulated Annealing

- Lets say there are 3 moves available, with changes in the objective function of  $d1 = -0.1$ ,  $d2 = 0.5$ ,  $d3 = -5$ . (Let  $T = 1$ ).
- pick a move randomly:
  - if  $d2$  is picked, move there.
  - if  $d1$  or  $d3$  are picked, probability of move =  $\exp(d/T)$
  - move 1:  $\text{prob1} = \exp(-0.1) = 0.9$ ,
    - i.e., 90% of the time we will accept this move
  - move 3:  $\text{prob3} = \exp(-5) = 0.05$ 
    - i.e., 5% of the time we will accept this move
- $T$  = “temperature” parameter
  - high  $T \Rightarrow$  probability of “locally bad” move is higher
  - low  $T \Rightarrow$  probability of “locally bad” move is lower
  - typically,  $T$  is decreased as the algorithm runs longer
    - i.e., there is a “temperature schedule”



# Simulated Annealing in Practice

Method proposed in 1983 by IBM researchers for solving VLSI layout problems (Kirkpatrick et al, *Science*, 220:671-680, 1983).

- theoretically will always find the global optimum (the best solution)


Useful for some problems, but can be very slow

- Slowness comes about because  $T$  must be decreased very gradually to retain optimality
- In practice how do we decide the rate at which to decrease  $T$ ? (this is a practical problem with this method)

# Genetic Algorithms (GA)

- Different approach to other search algorithms
  - A successor state is generated by combining two parent states rather than modifying the current state (**sexual reproduction rather than asexual reproduction**)
- A state is represented as a string over a finite alphabet (e.g. most commonly **0s** and **1s**, binary)
  - 8-queens
    - State = position of 8 queens each in a column  
=>  $8 \times \log(8)$  bits = 24 bits (for binary representation)
    - Alternatively, the state could be represented as digits, ranges 1 to 8
- Like Beam Search, GA starts with  $k$  randomly generated states (called the **population**)
- Objective function (**fitness function**).
  - Should return higher values for better states.
  - Opposite to heuristic function, e.g., # non-attacking pairs in 8-queens
- Produce the next generation of states by “simulated evolution”
  - Random selection
  - Crossover
  - Random mutation

# Genetic algorithms



(a)  
Initial Population

(b)  
Fitness Function

(c)  
Selection

(d)  
Cross-Over

(e)  
Mutation

- $23/(24+23+20+11) = 29\%$  etc


4 states for  
8-queens  
problem

2 pairs of 2 states randomly  
selected based on fitness.  
Random crossover points  
Selected using **Culling**

New states  
after crossover

Random  
mutation  
applied

# Genetic algorithms



Has the effect of “jumping” to a completely different new part of the search space (quite non-local)

# Genetic algorithm pseudocode

**function** GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

**inputs:** *population*, a set of individuals

          FITNESS-FN, a function that measures the fitness of an individual

**repeat**

*new\_population*  $\leftarrow$  empty set

**for**  $i = 1$  **to** SIZE(*population*) **do**

$x \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

$y \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

$child \leftarrow$  REPRODUCE( $x, y$ )

**if** (small random probability) **then**  $child \leftarrow$  MUTATE( $child$ )

    add  $child$  to *new\_population*

*population*  $\leftarrow$  *new\_population*

**until** some individual is fit enough, or enough time has elapsed

**return** the best individual in *population*, according to FITNESS-FN

---

**function** REPRODUCE( $x, y$ ) **returns** an individual

**inputs:**  $x, y$ , parent individuals

$n \leftarrow$  LENGTH( $x$ );  $c \leftarrow$  random number from 1 to  $n$

**return** APPEND(SUBSTRING( $x, 1, c$ ), SUBSTRING( $y, c + 1, n$ ))

# Comments on genetic algorithms

- Positive points
  - Random exploration can find solutions that local search can't
    - (via crossover primarily)
  - Appealing connection to human evolution
    - E.g., see related area of genetic programming, **Bioinformatics**
- Negative points
  - Large number of “tunable” parameters
    - Difficult to replicate performance from one problem to another
  - Lack of good empirical studies comparing to simpler methods
  - Useful on some (small?) set of problems but no convincing evidence that GAs are better than hill-climbing w/random restarts in general